

Relational Operations on MR and Hive

Relational Operations using MapReduce

1. Selection
 - a. Map : Read each row t of the table and check if it satisfies condition C , and if so, output (t,t)
 - b. Reduce : Do nothing
2. Projection
 - a. Map : Read each row t of the table and calculate subset of attributes t' and output (t,t')
 - b. Reduce : Eliminate duplicates
3. Union
 - a. Map : Reads in one file and send tuples as (t,t)
 - i. If multiple files exist, combine into one, but this combination is subjective
 - ii. Files may not have the same structure
 - iii. So each mapper reads in one input file and writes a key-value pair
 - b. Reducer : Merge the mapper outputs by eliminating duplicates
 - c. Schedule another map reduce job for selection
4. Intersection
 - a. Map : Read one file each and send tuples as (t,t)
 - b. Reducer : Consider only one duplicate per set of duplicates
5. Difference
 - a. Map : Read one file each and send tuples as (t,t)
 - b. Reducer : Consider non duplicates from mapper 1
6. Natural Join on attribute set B
 - a. Map : Each mapper reads $(B, (R,A))$
 - b. Reducer : Output (A,B,C)
7. Grouping and Aggregation
 - a. Map : For each line, output the attributes needed
 - b. Reduce : Aggregate each line

Hive

Hive is a system for querying and managing structured data built on top of MR and Hadoop. It has a SQL-like syntax for querying unstructured data as structured data, based on MR workflows.

The main difference between Hive and traditional RDBMS systems is that Hive performs schema on read, where it checks if the data conforms to the schema only when a query is issued. This makes for a very fast initial load and provides flexibility, but compromises on query time performance. Hence, Hive is used when schema is not known at load time.

Hive also does not support in place file updates, and hence produces delta files for each update, which are periodically written to the HDFS by MR jobs in the background, provisioned by the metastore.

1. Create -
 - a. CREATE TABLE table1 (count INT, bar STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LINES TERMINATED BY "\n" STORED AS TEXTFILE
2. Load -
 - a. LOAD DATA INPATH <hdfs file> INTO TABLE <table name>
 - b. Loads data without partitioning it
 - c. Keyword PARTITION can be used to partition this data
 - d. Makes access more efficient
 - e. Can specify LOCAL INPATH for local files
3. Select -
 - a. SELECT [ALL | DISTINCT] select_expr,... FROM <table> WHERE <cond> GROUP BY <cols> HAVING <cond> CLUSTER BY <cols> | SORT BY <cols> | DISTRIBUTE BY <cols> LIMIT <number>
 - b. Cluster by, sort by and distribute by all sort the data
 - c. Limit specifies the number of records to retrieve
4. Insert -
 - a. INSERT INTO TABLE <table> <select statement>
 - b. INSERT OVERWRITE TABLE <table> <select statement>

Hive uses a Hive QL Process Engine and an Execution engine running MR jobs based on the query, and uses a metastore to store metadata about the tables contained by the HDFS.

It's key components are

1. Hive CLI
2. Metastore
3. Compile - Translates statements into a DAG of MR jobs
4. Execution Engine
 - a. SerDe - Serialisation and Deserialisation is used.

The compiler is divided into the following components

1. Parser
 - a. Converts the query into a parse tree
2. Semantic Analyzer
 - a. Adds semantic information
 - b. Checks schema
 - c. Checks type of fields
3. Logical Plan Generator
 - a. Converts operators to logical operators
4. Optimizer
 - a. Optimizes the plan
5. Physical plan Generator
 - a. Converts plan to DAG of mapreduce jobs